

An answer is not a conclusion.

Same synthetic records. Same deterministic mock model. Different workflow. The one-call fixtures return a plausible candidate; THREADLINE preserves the rival contradiction and withholds aggregate release.

ANSWER FIRST

The baseline produced an answer. THREADLINE produced an auditable boundary around what the evidence can support.

What this packet proves

Evidence track	What is measured	What is not claimed
Deterministic 21-case replay	Fixed workflow behavior with identical case input	Live-model quality or field safety
Archived Prompt V2 live run	One provider extraction run plus deterministic policy	Same-model baseline comparison or repeated-run stability

Safety objective

False merges are treated as the central safety concern. Missed possible connections, dangerous unsupported connections, and correctly withheld cases remain distinct outcomes rather than being collapsed into a single flattering score.

Artifact THREADLINE-SUBMISSION-EVIDENCE-V1 / benchmark BENCH-4D40934964D7 / synthetic cases only / working tree dirty at generation: True

CASE / SAME INPUT

One plausible connection. One material rival.

Case INCIDENT-V1-RIVAL supplies the same 3 records to every compared system. Shared input SHA-256: bc8084e71ab87554163279a3b61d27d1b71f0c06fb3347c4555de5670ab83576.

Source record	Exact synthetic source text
FAMILY-018 family_report	Family report for Youssef Al Hassan, age 14. Speaks Arabic and some French. Last seen near Al Noor School at approximately 4:30 p.m. wearing a dark blue jacket. Small scar above the left eyebrow.
SHELTER-204 shelter_record	Shelter intake: Youssef Al Hassan, age 14. Arabic and limited French speaker. Arrived through North Gate at 7:10 p.m. wearing a navy blue jacket. Small scar above the left eyebrow.
HOSPITAL-052 hospital_intake	Hospital intake: Youssef Al Hassan, documented age 24 from a presented identity card. Admitted at South Clinic at 6:50 p.m. wearing a white shirt. No eyebrow scar recorded.

Observed deterministic outputs

System	Candidate records	Structured outcome
Generic single-prompt baseline	FAMILY-018, SHELTER-204	possible_candidate
Structured single-call baseline	FAMILY-018, SHELTER-204	possible_candidate
Full THREADLINE workflow	FAMILY-018, SHELTER-204	strong_candidate_for_review; aggregate withheld

SAFE CONCLUSION

The deterministic one-call fixtures produced a plausible review candidate. The full workflow preserved the rival age contradiction and withheld aggregate release. This demonstrates code-path behavior, not model quality or identity.

DETERMINISTIC COMPARISON

Coverage and restraint remain visible together.

Deterministic mock replay - not model performance.

System	Candidate recall	False-link rate	Correct abstention
Exact/fuzzy matching	4/12 (33.333%)	2/8 (25.000%)	0/1 (0.000%)
Generic single-prompt baseline	12/12 (100.000%)	8/8 (100.000%)	0/1 (0.000%)
Structured single-call baseline	12/12 (100.000%)	8/8 (100.000%)	0/1 (0.000%)
Full THREADLINE workflow	4/12 (33.333%)	0/8 (0.000%)	1/1 (100.000%)

Interpretation

The generic and structured one-call fixtures returned a positive review candidate in all eight different-identity cases. Full THREADLINE returned none and abstained on the one deliberately ambiguous case. These are fixture-and-rule outcomes, not estimates of live-model performance.

Ablations retained without spin

Configuration	Disabled	Primary or case-level effect
full	None	Changed metrics: None; affected cases: 0
no-prosecutor	prosecutor	Changed metrics: None; affected cases: 0
no-rivals	rivals	Changed metrics: None; affected cases: 0
no-adjudication	adjudicate	Changed metrics: None; affected cases: 0
selected-combined	adjudicate, prosecutor, rivals	Changed metrics: None; affected cases: 0

SOURCE-REMOVAL COUNTERFACTUAL

Removing HOSPITAL-052 changed aggregate contract status from withheld to released. This changes what the contract permits; it does not prove identity.

Evidence-contract-node removal is Not measured because the contract is not an independently disableable benchmark node in the current implementation.

ARCHIVED LIVE-PROVIDER EVIDENCE

One measured extraction run, reported separately.

Archived measured live-provider extraction run (one run).

Measure	Archived Prompt V2 result
Provider / model	openai_compatible / Qwen/Qwen3-30B-A3B-Instruct-2507
Prompt / temperature	v2 / 0.0
Successful extractions	57/58
Extraction TP / FP / FN	157 / 30 / 25
Extraction precision / recall / F1	0.8396 / 0.8626 / 0.8509
Candidate retrieval	14/14
Released proposals / false merges / false non-links	4 / 0 of 8 different-identity pairs / 10
Median latency	49,714 ms across 57 successful calls
p95 latency	82,819 ms; nearest-rank (ceil(0.95 * n))
Tokens / attempts	236,788 / 128
Approximate API cost	Not measured

What remains unknown

- One archived live-provider run; no repeated-run confidence interval.
- This evaluates live extraction followed by deterministic policy, not a same-model single-prompt baseline comparison.
- Synthetic records only; no operational humanitarian validation.
- Approximate cost is not reconstructed because the archive does not pin separate input/output token counts and a dated provider price table.
- The archived run did not record an exact git commit or dirty-tree patch; the current commit is not retroactively assigned.

Reproduce and inspect

Offline judge path: scripts/run-v1-demo.ps1
Submission evidence: cd backend; ../venv/Scripts/python.exe scripts/build_submission_evidence.py
Auditable sources: results.json, results.csv, raw-outputs.json, ablation-counterfactual.json, benchmark-report.md

FINAL BOUNDARY

All cases and identities are fictional. This research prototype is not validated for operational humanitarian use and cannot establish a person's identity.